

# CodeNotebook

Scott McCoy - U80152879

2/27/2021

## Mobile Price Classification

### Models

- Linear Models
  - Nnet Multiclass Linear Model (Baseline)
  - Binary Logistic Regression
  - Multiple Binary Logistic Regressions
  - Layered Binary Logistic Regressions
- Tree-based Models
  - Decision Tree
  - Random Forest
    - Multiple Binary Random Forests
    - Multiclass Random Forest

## Linear Models

### Neural Net Multiclass Linear Model

There are a number of methods that can be applied to a multiclass classification problem.

Of the various models we applied to our problem, the best performing was a log-linear neural network model.

We used this as a baseline to inform our strategy and compare results to the methods discusses in the course.

```
defaultW <- getOption("warn")
options(warn = -1)
options(warn = defaultW)

set.seed(430)
library(data.table)
library(nnet)
mo <- fread('~\\R\\mobile\\train.csv')
mo_obs <- nrow(mo)
mo_idx <- sample(mo_obs, size = trunc(0.70 * mo_obs))
mo_trn <- mo[mo_idx, ]
mo_test <- mo[-mo_idx, ]

Y_train <- mo_trn[,price_range]
Y_test <- mo_test[,price_range]

model_multi <- multinom(price_range ~ ., data = mo_trn) # instantiating model
```

```
## # weights:  88 (63 variable)
## initial  value 1940.812106
## iter   10 value 1539.449569
## iter   20 value 1372.517934
## iter   30 value 1328.431365
## iter   40 value 1203.483614
## iter   50 value  826.577171
## iter   60 value  446.763231
## iter   70 value  59.572210
## iter   80 value  31.622093
## iter   90 value  21.277743
## iter  100 value  15.526852
## final   value  15.526852
## stopped after 100 iterations
```

```
# train predictions
Y_train_hat_df <- predict(model_multi, newdata = mo_trn, type = "prob")

Y_train_hat <- data.table(colnames(Y_train_hat_df)[max.col(Y_train_hat_df,ties.method="first")])
Y_train_hat <- transform(Y_train_hat, V1 = as.numeric(V1))

train_accuracy <- mean(Y_train == Y_train_hat)

#test predictions
Y_test_hat_df <- predict(model_multi, newdata = mo_test, type = "prob")

Y_test_hat <- data.table(colnames(Y_test_hat_df)[max.col(Y_test_hat_df,ties.method="first")])
Y_test_hat <- transform(Y_test_hat, V1 = as.numeric(V1))[,V1]

test_accuracy <- mean(Y_test == Y_test_hat)

# Evaluation:
cm <- table(observed=Y_test, predicted=Y_test_hat) # confusion matrix

diag = diag(cm) # number of correctly classified instances per class
rowsums = apply(cm, 1, sum) # number of instances per class
colsums = apply(cm, 2, sum) # number of predictions per class

precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)

scores <- data.frame(precision, recall, f1)

train_accuracy
```

```
## [1] 0.9971429
```

```
test_accuracy
```

```
## [1] 0.9716667
```

```
cm
```

```
##          predicted
## observed  0    1    2    3
##          0 147    1    0    0
##          1   2 147    2    0
##          2   0   5 147    6
##          3   0   0   1 142
```

```
scores
```

```
## precision    recall      f1
## 0 0.9865772 0.9932432 0.9898990
## 1 0.9607843 0.9735099 0.9671053
## 2 0.9800000 0.9303797 0.9545455
## 3 0.9594595 0.9930070 0.9759450
```

We have a test accuracy of 97% for this model. This high accuracy of this linear model indicates that there is likely a strong linear relationship between the predictors and price\_range.

## Binary Logistic Regression

Because we have more than 2 classes of predictors, we cannot simply predict between all 4 with a simple logistic regression.

Our first approach was to predict between the the lowest price range (0,1) and the highest price (2,3) with a binomial model.

```

defaultW <- getOption("warn")
options(warn = -1)

library(caTools)
library(caret)
mo_p <- fread('~R/mobile/processed_train.csv')
m_binary <- mo_p[, !c("price_range", "p0", "p1", "p2", "p3")]
sampleSplit <- sample.split(Y=m_binary$price_binary, SplitRatio=0.7)
trainSet <- subset(x=m_binary, sampleSplit==TRUE)
testSet <- subset(x=m_binary, sampleSplit==FALSE)

log_model_binary <- glm(price_binary ~ ., family=binomial(link='logit'), data=trainSet)

probabs <- predict(log_model_binary, testSet[,!c("price_binary")],type='response')
preds <- ifelse(probabs > 0.5, 1, 0)

test_accuracy <- mean(testSet$price_binary == preds)

cm <- table(observed=testSet$price_binary, predicted=preds)

test_accuracy

```

```
## [1] 0.995
```

```
cm
```

```

##      predicted
## observed  0   1
##      0 300   0
##      1   3 297

```

This model does an excellent job differentiating between low prices and high prices.

## Multiple Binary Logistic Regression

Our next step was to run multiple binomial models, one for each class of price range.

We can select between all 4 classes by ultimately predicting the class with the highest probability between all 4 models.

```
mo_obs <- nrow(mo_p)
mo_idx <- sample(mo_obs, size = trunc(0.70 * mo_obs))
mo_trn <- mo_p[mo_idx, ]
mo_test <- mo_p[-mo_idx, ]

# 4 separate train and test sets
X_train <- mo_trn[,1:20]
Y_train <- mo_trn[,price_range]
Y_train0 <- mo_trn[,p0]
Y_train1 <- mo_trn[,p1]
Y_train2 <- mo_trn[,p2]
Y_train3 <- mo_trn[,p3]
X_train0 <- data.table(X_train)
X_train0[,p0 := Y_train0]
X_train1 <- data.table(X_train)
X_train1[,p1 := Y_train1]
X_train2 <- data.table(X_train)
X_train2[,p2 := Y_train2]
X_train3 <- data.table(X_train)
X_train3[,p3 := Y_train3]

X_test <- mo_test[,1:20]
Y_test <- mo_test[,price_range]
Y_test0 <- mo_test[,p0]
Y_test1 <- mo_test[,p1]
Y_test2 <- mo_test[,p2]
Y_test3 <- mo_test[,p3]
X_test0 <- data.table(X_test)
X_test0[,p0 := Y_test0]
X_test1 <- data.table(X_test)
X_test1[,p1 := Y_test1]
X_test2 <- data.table(X_test)
X_test2[,p2 := Y_test2]
X_test3 <- data.table(X_test)
X_test3[,p3 := Y_test3]

# fitting models
glm.fit0 <- glm(p0 ~ ., data = X_train0, family = binomial)
glm.fit1 <- glm(p1 ~ ., data = X_train1, family = binomial)
glm.fit2 <- glm(p2 ~ ., data = X_train2, family = binomial)
glm.fit3 <- glm(p3 ~ ., data = X_train3, family = binomial) # binomial for logistic regression

# train predictions
Y_train_hat0 <- predict(glm.fit0, newdata = X_train0, type = "response")
Y_train_hat1 <- predict(glm.fit1, newdata = X_train1, type = "response")
Y_train_hat2 <- predict(glm.fit2, newdata = X_train2, type = "response")
Y_train_hat3 <- predict(glm.fit3, newdata = X_train3, type = "response")

Y_train_hat_df <- data.table("0" = Y_train_hat0, "1" = Y_train_hat1, "2" = Y_train_hat2, "3" = Y_train_hat3)

Y_train_hat <- data.table(colnames(Y_train_hat_df)[max.col(Y_train_hat_df, ties.method="first")])
Y_train_hat <- lapply(Y_train_hat[,], as.numeric)
```

```

train_accuracy <- mean(Y_train == Y_train_hat$V1)

#test predictions
Y_test_hat0 <- predict(glm.fit0, newdata = X_test0, type = "response")
Y_test_hat1 <- predict(glm.fit1, newdata = X_test1, type = "response")
Y_test_hat2 <- predict(glm.fit2, newdata = X_test2, type = "response")
Y_test_hat3 <- predict(glm.fit3, newdata = X_test3, type = "response")

Y_test_hat_df <- data.table("0" = Y_test_hat0, "1" = Y_test_hat1, '2' = Y_test_hat2, "3" = Y_test_hat3)

Y_test_hat <- data.table(colnames(Y_test_hat_df)[max.col(Y_test_hat_df,ties.method="first")])
Y_test_hat <- lapply(Y_test_hat[,], as.numeric)

test_accuracy <- mean(Y_test == Y_test_hat$V1)

# Evaluation
cm <- table(observed=Y_test, predicted=Y_test_hat$V1)

diag = diag(cm) # number of correctly classified instances per class
rowsums = apply(cm, 1, sum) # number of instances per class
colsums = apply(cm, 2, sum) # number of predictions per class

precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)

scores <- data.frame(precision, recall, f1)
train_accuracy

```

```
## [1] 0.8857143
```

```
test_accuracy
```

```
## [1] 0.8516667
```

```
cm
```

```
##      predicted
## observed  0   1   2   3
##      0 170   2   1   0
##      1   2 100  45   0
##      2   0  31 107   3
##      3   0   0   5 134
```

```
scores # has trouble distinguishing between 1 and 2
```

```
## precision recall f1
## 0 0.9883721 0.9826590 0.9855072
## 1 0.7518797 0.6802721 0.7142857
## 2 0.6772152 0.7588652 0.7157191
## 3 0.9781022 0.9640288 0.9710145
```

This set of models identifies the lowest and highest price ranges well, but has trouble distinguishing between price range 1 and 2.

## Layered Binary Regressions

Our binary first model predicts well between low and high prices (0,1) vs (2,3), and our multiple binary models do a good job predicting the lowest (0) and highest (3) price ranges.

We decided to combine the strengths of both sets of models by initially predicting high vs low, then predicting 0 vs 1 for low predictions, or 2 vs 3 for high predictions.

```

Y_train <- mo_trn[,price_binary]
X_train <- mo_trn[, 1:20]
X_train[,price_binary := Y_train] # training set - first binary (low-high) model
L
X_train01 <- mo_trn[price_range < 2, c(1:20,24) ] # training set - second layer binary model
(0,1)
X_train23 <- mo_trn[price_range > 1, c(1:20,26) ] # training set - second layer binary model
(2,3)

Y_test <- mo_test[,price_binary]
X_test <- mo_test[, 1:20]
X_test[, price_binary := Y_test]

# fitting models
glm.fit <- glm(price_binary ~ ., data = X_train, family = binomial) # binomial for logistic regression

glm.fit01 <- glm(p1 ~ ., data = X_train01, family = binomial)
glm.fit23 <- glm(p3 ~ ., data = X_train23, family = binomial)

#####
# 1st Prediction Layer

#test predictions
Y_test_hat <- predict(glm.fit, newdata = X_test, type = "response")
Y_test_hat <- data.table(Y_test_hat > 0.5)
Y_test_hat <- transform(Y_test_hat, V1 = as.numeric(V1))[,V1]

binary1_accuracy <- mean(Y_test == Y_test_hat)

x2 <- data.table(mo_test)
x2[, hi_lo_prediction := Y_test_hat]

low_table <- x2[hi_lo_prediction == 0, c(1:20, 24, 21)]
high_table <- x2[hi_lo_prediction == 1, c(1:20, 26, 21)] # split based on prediction from first layer

#####
# 2nd Prediction Layer

# if first layer predicted low, second layer predicts 0 or 1
# if first layer predicted high, second layer predicts 2 or 3

#test predictions
Y_test_hat_low <- predict(glm.fit01, newdata = low_table[, 1:21], type = "response")
Y_test_hat_low <- data.table(Y_test_hat_low > 0.5)
Y_test_hat_low <- transform(Y_test_hat_low, V1 = as.numeric(V1))[,V1]

test_accuracy_low <- mean(low_table$price_range == Y_test_hat_low) # price_range column - not used as input, just for checking results

Y_test_hat_high <- predict(glm.fit23, newdata = high_table[, 1:21], type = "response")

```



```

Y_test_hat_high <- data.table(Y_test_hat_high > 0.5)
Y_test_hat_high <- Y_test_hat_high + 2
Y_test_hat_high <- transform(Y_test_hat_high, V1 = as.numeric(V1))[,V1]

test_accuracy_high <- mean(high_table$price_range == Y_test_hat_high) # price_range column - not used as input, just for checking results

# Combining Results
matrix1 <- data.table(low_table)
matrix1[, prediction := Y_test_hat_low]
matrix1[, p1 := NULL]
matrix2 <- data.table(high_table)
matrix2[, prediction := Y_test_hat_high]
matrix2[, p3 := NULL]

results <- rbind(matrix1, matrix2)

# Evaluation

test_accuracy <- mean(results$price_range == results$prediction)

cm <- table(observed=results$price_range, predicted=results$prediction)

diag = diag(cm) # number of correctly classified instances per class
rowsums = apply(cm, 1, sum) # number of instances per class
colsums = apply(cm, 2, sum) # number of predictions per class

precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)

scores <- data.frame(precision, recall, f1)

binary1_accuracy

```

```
## [1] 0.9766667
```

```
test_accuracy_low
```

```
## [1] 0.9713376
```

```
test_accuracy_high
```

```
## [1] 0.9370629
```

```
test_accuracy #overall accuracy for final output through both layers
```

```
## [1] 0.955
```

```
cm
```

```
##          predicted
## observed  0    1    2    3
##          0 170    3    0    0
##          1  2 135   10    0
##          2  0  4 134    3
##          3  0  0  5 134
```

```
scores
```

```
## precision    recall      f1
## 0 0.9883721 0.9826590 0.9855072
## 1 0.9507042 0.9183673 0.9342561
## 2 0.8993289 0.9503546 0.9241379
## 3 0.9781022 0.9640288 0.9710145
```

This approach yields accuracy almost as high as our baseline neural net model.

Logistic Models Summary: \* Having more than 2 predictors complicated problem, and we had to be creative with models we applied to this dataset. \* These models overall had high accuracy, indicating a linear relationship between predictors and price range.

## Tree-based Models

Tree Models have the advantage of being easily adaptable to a multiclass classification problem, and are easily interpretable.

We apply several multiclass and multiple binary tree models below.

### Simple Decision Tree

```
library(rpart)
library(rpart.plot)

split = sample.split(mo$price_range, SplitRatio = 0.7)
data_train = subset(mo, split == TRUE)
data_test = subset(mo, split == FALSE)

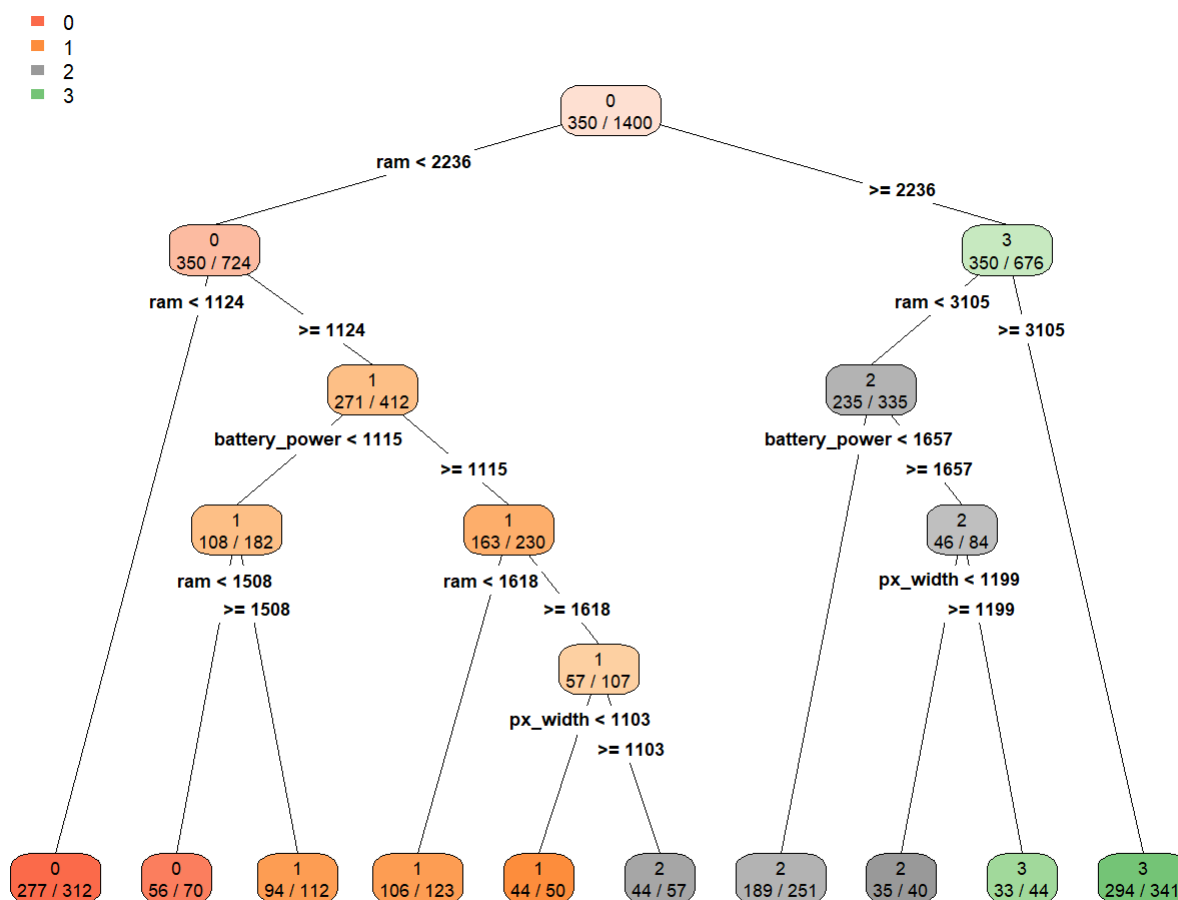
tree = rpart(price_range ~ ., method = "class", data = data_train)

tree.pred = predict(tree, newdata = data_test, type = 'class')
tree.accuracy = mean(tree.pred == data_test$price_range)
tree.accuracy
```

```
## [1] 0.8
```

## # Pruning Tree

```
tree2 <- prune(tree, cp = 0.01000000)
rpart.plot(tree2, type = 4, branch = 0, extra = 2)
```



```
CFit1 <- predict(tree2, data_test, type = "class")
#ConfM1 <- table(data_train$price_range, CFit1)
#(E1 <- (sum(ConfM1) - sum(diag(ConfM1)))/sum(ConfM1))
tree.accuracy = mean(CFit1 == data_test$price_range)
tree.accuracy
```

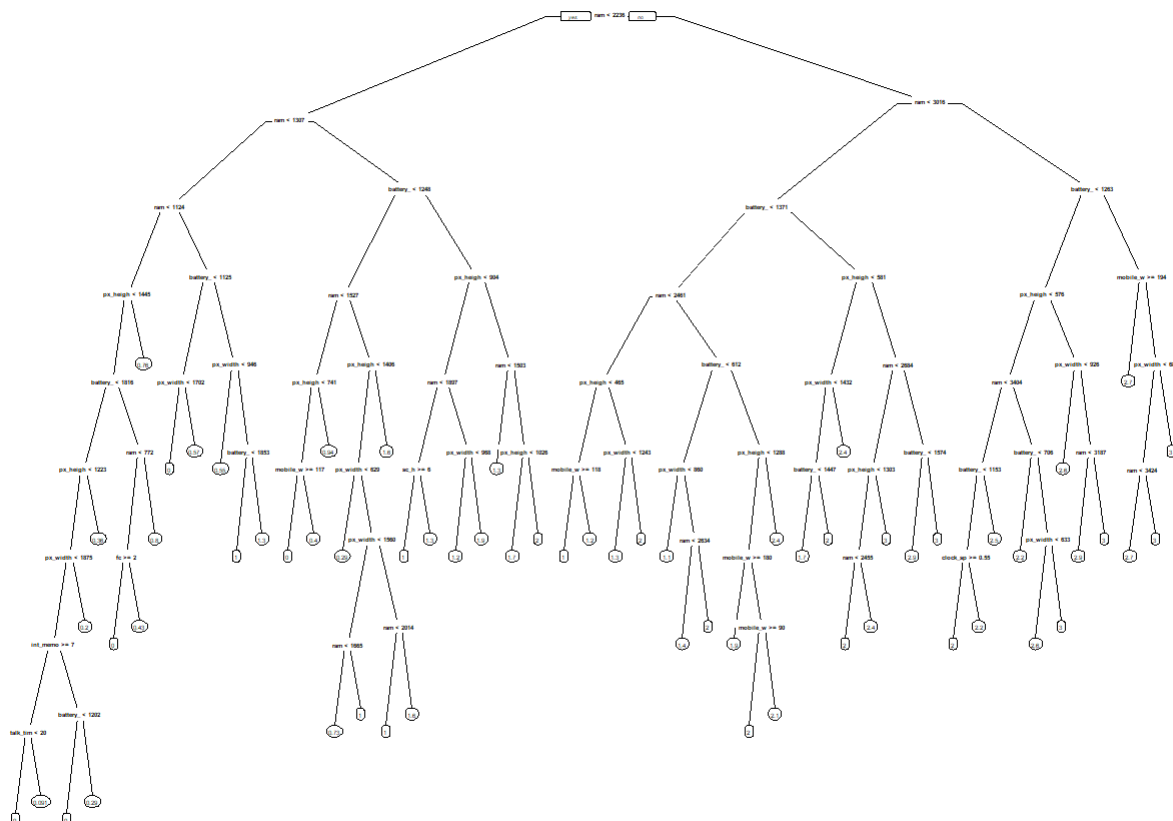
```
## [1] 0.8
```

## # Cross Validation

```
tr.control = trainControl(method = "cv", number = 10)
cp.grid = expand.grid(cp = (0:10)*0.001)
tr = train(price_range ~., data = data_train, method = "rpart", trControl = tr.control, tuneGrid = cp.grid)
tr
```

```
## CART
##
## 1400 samples
## 20 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1260, 1260, 1260, 1260, 1260, 1260, ...
## Resampling results across tuning parameters:
##
##   cp      RMSE      Rsquared   MAE
## 0.000 0.3555902 0.8992706 0.1818983
## 0.001 0.3752021 0.8874764 0.1981023
## 0.002 0.3806766 0.8836119 0.2117234
## 0.003 0.3851284 0.8810468 0.2218490
## 0.004 0.3973982 0.8739792 0.2412526
## 0.005 0.4220260 0.8576575 0.2890290
## 0.006 0.4292327 0.8526930 0.3074388
## 0.007 0.4325203 0.8501444 0.3164713
## 0.008 0.4356231 0.8479527 0.3232046
## 0.009 0.4382216 0.8460195 0.3305465
## 0.010 0.4428475 0.8425110 0.3235908
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was cp = 0.
```

```
# Predictions with best tree from CV
best.tree = tr$finalModel
prp(best.tree)
```



```
best.tree.pred = predict(best.tree, newdata = data_test)
tree.accuracy = mean(best.tree.pred == data_test$price_range)
tree.accuracy
```

```
## [1] 0.6216667
```

## Random Forests

### Multiple Binary Random Forest Classifiers

```
library(mltools)
library(randomForest)
library(tidyr)
library(tidyselect)
library(ggplot2)
library(dplyr)

mobile_data <- fread('~R/mobile/train.csv', stringsAsFactors = T)

#set price_range as factor for one hot encoding
mobile_data$price_range <- as.factor(mobile_data$price_range)

#one hot encode the data for price range
mobile_data_one = one_hot(mobile_data,cols='price_range')

#split the data into training and test
mobile_data_one[, test:=0]
mobile_data_one[sample(nrow(mobile_data_one), 300), test:=1] # take 300 random rows and stick them in the test set
# now split
mobile_data_one_test <- mobile_data_one[test==1]
mobile_data_one_train <- mobile_data_one[test==0]

### Train data for each price level randomForest model, setting target variable as a factor
mobile_train_0 <- mobile_data_one_train %>% select(-(price_range_1:test))
mobile_train_0$price_range_0 <- as.factor(mobile_train_0$price_range_0)

mobile_train_1 <- mobile_data_one_train %>% select(-c(price_range_0,(price_range_2:test)))
mobile_train_1$price_range_1 <- as.factor(mobile_train_1$price_range_1)

mobile_train_2 <- mobile_data_one_train %>% select(-c((price_range_0:price_range_1),(price_range_3:test)))
mobile_train_2$price_range_2 <- as.factor(mobile_train_2$price_range_2)

mobile_train_3 <- mobile_data_one_train %>% select(-c((price_range_0:price_range_2),test))
mobile_train_3$price_range_3 <- as.factor(mobile_train_3$price_range_3)

###Test data
mobile_predictors_test <- mobile_data_one_test %>% select(-(price_range_0:test))
#instantiate test Ys
price_0_test <- mobile_data_one_test %>% select(price_range_0)
price_0_test_f <- as.factor(price_0_test$price_range_0)

price_1_test <- mobile_data_one_test %>% select(price_range_1)
price_1_test_f <- as.factor(price_1_test$price_range_1)

price_2_test <- mobile_data_one_test %>% select(price_range_2)
price_2_test_f <- as.factor(price_2_test$price_range_2)

price_3_test <- mobile_data_one_test %>% select(price_range_3)
price_3_test_f <- as.factor(price_3_test$price_range_3)
```

#####

```

####Code below did not work for the random Forest model but
####could be used for other applications

#cross validation (?)

#separate X (predictors)
mobile_predictors_train <- mobile_data_one_train %>% select(-(price_range_0:test))

#instantiate each individual train Ys and obtain the vector of the values
price_0_train <- mobile_data_one_train %>% select(price_range_0)
y_0_train <- price_0_train$price_range_0

price_1_train <- mobile_data_one_train %>% select(price_range_1)
y_1_train <- price_1_train$price_range_1

price_2_train <- mobile_data_one_train %>% select(price_range_2)
y_2_train <- price_2_train$price_range_2

price_3_train <- mobile_data_one_train %>% select(price_range_3)
y_3_train <- price_3_train$price_range_3

#####
#fit the models for each price level
#Random Forest Classifier for price range 0
fit.rndfor_0 <- randomForest(price_range_0 ~.,
                             data = mobile_train_0,
                             importance = TRUE,
                             xtest = mobile_predictors_test,
                             ytest = price_0_test_f)
#Random Forest Classifier for price range 1
fit.rndfor_1 <- randomForest(price_range_1 ~.,
                             data = mobile_train_1,
                             importance=TRUE,
                             xtest = mobile_predictors_test,
                             ytest = price_1_test_f)
#Random Forest Classifier for price range 2
fit.rndfor_2 <- randomForest(price_range_2 ~.,
                             data = mobile_train_2,
                             importance=TRUE,
                             xtest = mobile_predictors_test,
                             ytest = price_2_test_f)
#Random Forest Classifier for price range 3
fit.rndfor_3 <- randomForest(price_range_3 ~.,
                             data = mobile_train_3,
                             importance=TRUE,
                             xtest = mobile_predictors_test,
                             ytest = price_3_test_f)

#Analyze the results
# Price Range 0 train
y_hat_0 <- fit.rndfor_0$predicted
price_0_acc <- mean(y_hat_0 == y_0_train)
# Price Range 0 test

```

```

y_test_hat_0 <- fit.rndfor_0$test$predicted
price_0_acc_test <- mean(y_test_hat_0 == price_0_test$price_range_0)

#Price Range 1 train
y_hat_1 <- fit.rndfor_1$predicted
price_1_acc <- mean(y_hat_1 == y_1_train)
# Price Range 1 test
y_test_hat_1 <- fit.rndfor_1$test$predicted
price_1_acc_test <- mean(y_test_hat_1 == price_1_test$price_range_1)

#Price Range 2 train
y_hat_2 <- fit.rndfor_2$predicted
price_2_acc <- mean(y_hat_2 == y_2_train)
# Price Range 2 test
y_test_hat_2 <- fit.rndfor_2$test$predicted
price_2_acc_test <- mean(y_test_hat_2 == price_2_test$price_range_2)

#Price Range 3 train
y_hat_3 <- fit.rndfor_3$predicted
price_3_acc <- mean(y_hat_3 == y_3_train)
# Price Range 3 test
y_test_hat_3 <- fit.rndfor_3$test$predicted
price_3_acc_test <- mean(y_test_hat_3 == price_3_test$price_range_3)

#####
#Code below has the purpose of combining all 4 previous models

#Building a model for a prediction with all models
#set up y_test as a 4 level factor
price_levels <- mobile_data_one_test %>% select((price_range_0:price_range_3))
colnames(price_levels) <- c("0","1","2","3")
w <- which(price_levels==1,arr.ind = T)
mobile_data_one_test$price_level <- toupper(names(price_levels)[w[order(w[,1]),2]])
#Add these values into a data table
prediction_dt <- data.table("0" = fit.rndfor_0$test$votes[,2],
                           "1" = fit.rndfor_1$test$votes[,2],
                           "2" = fit.rndfor_2$test$votes[,2],
                           "3" = fit.rndfor_3$test$votes[,2])

label_rf <- apply(prediction_dt,1,which.max)-1
decision_dt <- data.table("predicted values"= label_rf)

decision_dt$actualvalues <- mobile_data_one_test$price_level

#Evaluate decisions
y_test <- as.numeric(decision_dt$actualvalues)
predictions <- as.numeric(decision_dt$`predicted values`)
analysis_table <- table(y_test,predictions)

diag = diag(analysis_table) # number of correctly classified instances per class
rowsums = apply(analysis_table, 1, sum) # number of instances per class
colsums = apply(analysis_table, 2, sum) # number of predictions per class

```



```
precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)
```

```
scores <- data.frame(precision, recall, f1)
test_accuracy <- mean(y_test == predictions)
```

```
scores
```

```
##   precision    recall      f1
## 0 0.8961039 0.9324324 0.9139073
## 1 0.7662338 0.7468354 0.7564103
## 2 0.7625000 0.7530864 0.7577640
## 3 0.8939394 0.8939394 0.8939394
```

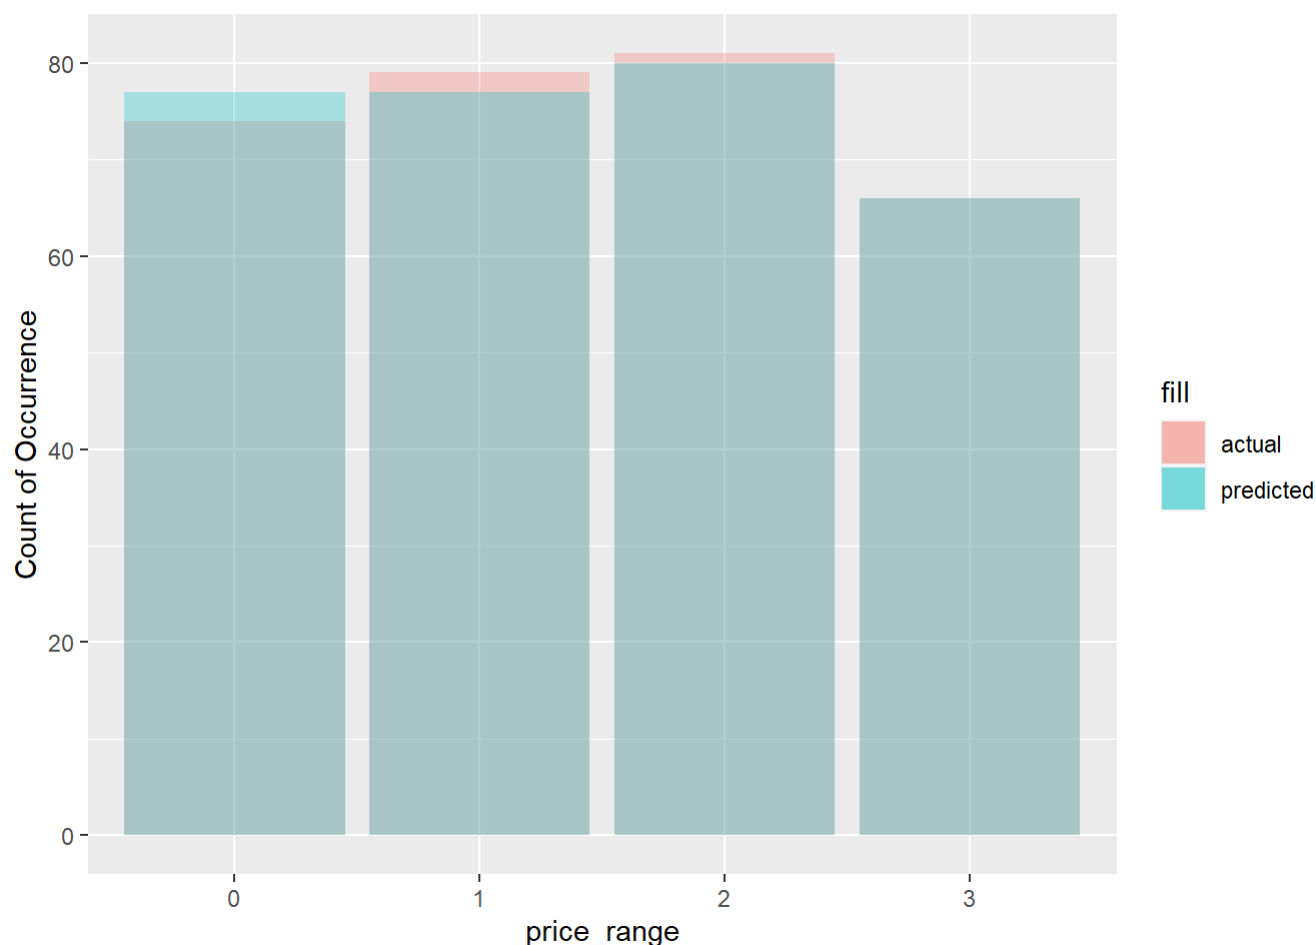
```
analysis_table
```

```
##      predictions
## y_test 0 1 2 3
##      0 69 5 0 0
##      1 8 59 12 0
##      2 0 13 61 7
##      3 0 0 7 59
```

```
test_accuracy
```

```
## [1] 0.8266667
```

```
#graphing the results
df_rf <- data.table(rowsums)
df_rf$pred <- colsums
df_rf$price_range <- c("0",'1','2','3')
colnames(df_rf) <- c('actual','predicted','price_range')
ggplot(NULL,aes(x=price_range,y=actual))+
  geom_bar(aes(fill="actual"), data= df_rf, stat = 'identity',position = "dodge",alpha=0.3)+
  geom_bar(aes(y=predicted,fill="predicted"),data=df_rf,stat = 'identity',position = "dodge",alp
ha = 0.3)+
  ylab("Count of Occurrence")
```



## Multiclass Random Forest Classifier

Here we use cross validation to find optimal values of 2 Random Forest hyperparameters \* m - the number of predictors randomly selected for each split \* number of trees

```
mo <- fread('~R/mobile/train.csv')
mo$price_range <- as.factor(mo$price_range)
mo_obs <- nrow(mo)
mo_idx <- sample(mo_obs, size = trunc(0.70 * mo_obs))
mo_trn <- mo[mo_idx, ]
mo_test <- mo[-mo_idx, ]

Y_test <- mo_test[,price_range]

#https://www.blopig.com/blog/2017/04/a-very-basic-introduction-to-random-forests-using-r/

dd <- data.table()
dd

## Null data.table (0 rows and 0 cols)
```

```

p <- length(colnames(mo_trn)) -1
p_ov2 <- p / 2
p_sqrt <- sqrt(length(colnames(mo_trn)) -1)

trees <- seq(from = 10, to = 210, by = 10)

for (num_pred in c(p, p_ov2, p_sqrt)) {
  CV_accuracies = c()
  for (num_trees in trees){

    #Perform K-fold cross validation
    k = 5
    #Randomly shuffle the data
    mo_trn_cross <- mo_trn[sample(nrow(mo_trn)),]

    #Create K equally size folds
    folds <- cut(seq(1,nrow(mo_trn_cross)),breaks=k,labels=FALSE)

    accuracies <- c()

    #Perform K-fold cross validation

    for(i in 1:k){
      #Segment data by fold with which() function
      testIndexes <- which(folds==i,arr.ind=TRUE)
      testData <- mo_trn_cross[testIndexes, ]
      trainData <- mo_trn_cross[-testIndexes, ]
      Y_CV <- testData$price_range

      #num_features <- sqrt(length(colnames(mo_trn)) -1)

      rf_classifier <- randomForest(price_range ~ ., data = trainData, ntree = num_trees, mtry =
num_pred, importance = TRUE )
      Y_test_hat <- predict(rf_classifier, newdata = testData, type = "class")

      accuracy <- mean(Y_test_hat == Y_CV)

      accuracies <- c(accuracies, accuracy)
    }
    CV_accuracy <- mean(accuracies)
    CV_accuracies <- c(CV_accuracies, CV_accuracy)
  }
  dd <- cbind(dd, CV_accuracies)

}
dd[, num_tree := trees]
dd

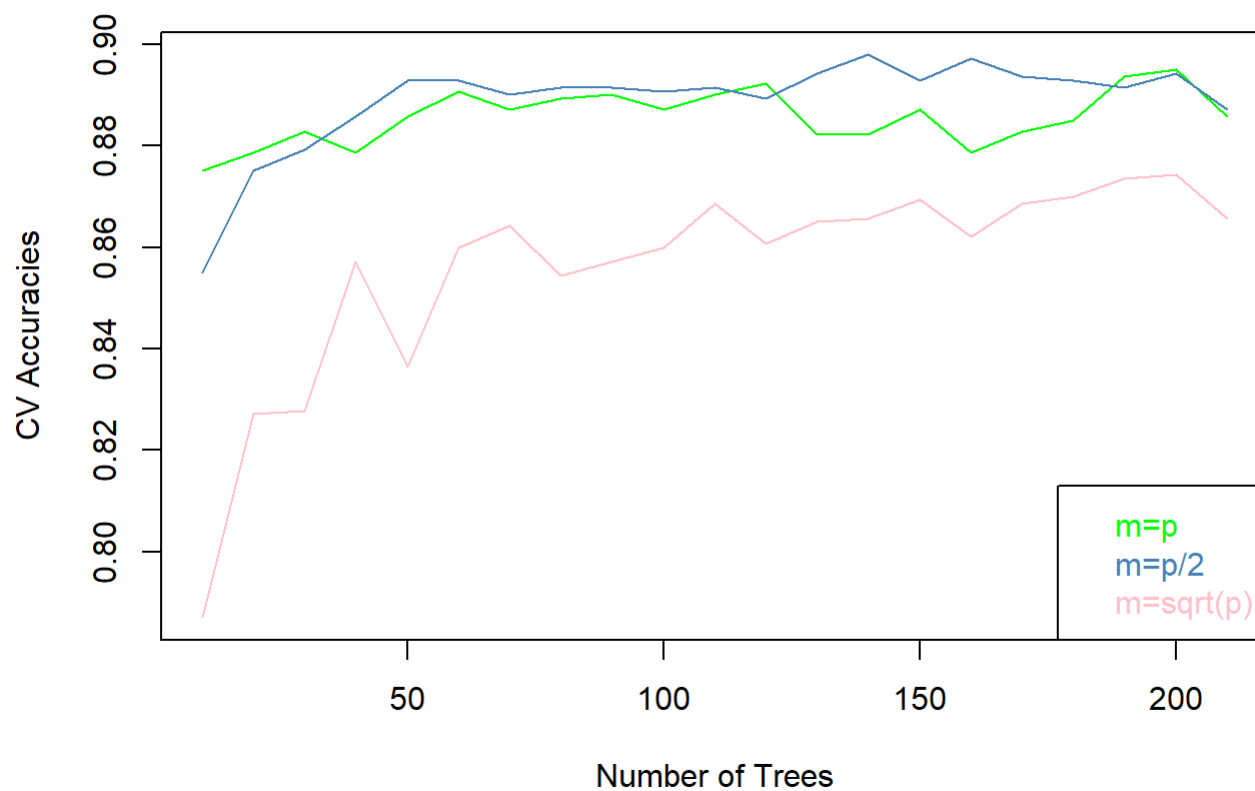
```

```
##      CV_accuracies CV_accuracies CV_accuracies num_tree
## 1:      0.8750000      0.8550000      0.7871429      10
## 2:      0.8785714      0.8750000      0.8271429      20
## 3:      0.8828571      0.8792857      0.8278571      30
## 4:      0.8785714      0.8857143      0.8571429      40
## 5:      0.8857143      0.8928571      0.8364286      50
## 6:      0.8907143      0.8928571      0.8600000      60
## 7:      0.8871429      0.8900000      0.8642857      70
## 8:      0.8892857      0.8914286      0.8542857      80
## 9:      0.8900000      0.8914286      0.8571429      90
## 10:      0.8871429      0.8907143      0.8600000     100
## 11:      0.8900000      0.8914286      0.8685714     110
## 12:      0.8921429      0.8892857      0.8607143     120
## 13:      0.8821429      0.8942857      0.8650000     130
## 14:      0.8821429      0.8978571      0.8657143     140
## 15:      0.8871429      0.8928571      0.8692857     150
## 16:      0.8785714      0.8971429      0.8621429     160
## 17:      0.8828571      0.8935714      0.8685714     170
## 18:      0.8850000      0.8928571      0.8700000     180
## 19:      0.8935714      0.8914286      0.8735714     190
## 20:      0.8950000      0.8942857      0.8742857     200
## 21:      0.8857143      0.8871429      0.8657143     210
##      CV_accuracies CV_accuracies CV_accuracies num_tree
```

```
colnames(dd) <- c("m=p", "m=p/2", "m=sqrt(p)", "num_tree")

plot(dd$m=p ~ dd$num_tree, type = 'l', ylim = c(min(dd[,1:3]), max(dd[,1:3])), col = 'green',
     xlab = 'Number of Trees', ylab = "CV Accuracies")
lines(dd$num_tree, dd$m=p/2, col = "steelblue")
lines(dd$num_tree, dd$m=sqrt(p), col = "pink")

legend("bottomright",
      legend = c("m=p", "m=p/2", "m=sqrt(p)"),
      text.col = c("green", "steelblue", "pink")
)
```



```
which.max(dd$m=p/2) # 140 trees with  $m = p/2$  is max cv accuracy
```

```
## [1] 14
```

```

rf_classifier <- randomForest(price_range ~ ., data = mo_trn, ntree = 140, mtry = p_ov2, importance = TRUE )

Y_test_hat <- predict(rf_classifier, newdata = mo_test, type = "class")

test_accuracy <- mean(Y_test_hat == Y_test)

cm <- table(observed=Y_test, predicted=Y_test_hat)

diag = diag(cm) # number of correctly classified instances per class
rowsums = apply(cm, 1, sum) # number of instances per class
colsums = apply(cm, 2, sum) # number of predictions per class

precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)

scores <- data.frame(precision, recall, f1)

test_accuracy

```

```
## [1] 0.8866667
```

```
cm
```

```

##      predicted
## observed  0   1   2   3
##      0 149  10   0   0
##      1  12 113   8   0
##      2   0  11 117  13
##      3   0   0  14 153

```

```
scores
```

```

## precision recall      f1
## 0 0.9254658 0.9371069 0.9312500
## 1 0.8432836 0.8496241 0.8464419
## 2 0.8417266 0.8297872 0.8357143
## 3 0.9216867 0.9161677 0.9189189

```

```
options(warn = defaultW)
```

Typically the recommended value for  $m$  in a random forest is the square root of the number of predictors, but that approach consistently underperformed compared to  $m = p$  (same as bagging), and  $m = p / 2$ .

The best cv accuracy came from a model with 140 trees and  $m = p / 2$ .