

RandomForest

Scott McCoy - U80152879

2/24/2021

Random Forest with K-Fold CV

Libraries:

```
set.seed(430)
library(data.table)
library(ggplot2)
library(ggthemes)
library(scales)
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##     margin
```

```
theme_set(theme_bw())
```

Train-Test Split

```
mo <- fread('~R/mobile/train.csv')
mo$price_range <- as.factor(mo$price_range)
mo_obs <- nrow(mo)
mo_idx <- sample(mo_obs, size = trunc(0.70 * mo_obs))
mo_trn <- mo[mo_idx, ]
mo_test <- mo[-mo_idx, ]

Y_test <- mo_test[,price_range]
```

Finding optimal hyperparameter values with K-Fold CV:

```

dd <- data.table() # empty data table to fill with cv accuracies over loop

p <- length(colnames(mo_trn)) -1 # variables for number of predictors hyperparameter
p_ov2 <- p / 2
p_sqrt <- sqrt(length(colnames(mo_trn)) -1)

trees <- seq(from = 10, to = 210, by = 10) # sequence of # of trees hyperparameter

for (num_pred in c(p, p_ov2, p_sqrt)) {
  CV_accuracies = c()
  for (num_trees in trees){

    #Perform K-fold cross validation
    k = 5
    #Randomly shuffle the data
    mo_trn_cross <- mo_trn[sample(nrow(mo_trn)),]

    #Create K equally size folds
    folds <- cut(seq(1,nrow(mo_trn_cross)),breaks=k,labels=FALSE)

    accuracies <- c()

    #Perform K-fold cross validation

    for(i in 1:k){
      #Segment your data by fold using the which() function
      testIndexes <- which(folds==i,arr.ind=TRUE)
      testData <- mo_trn_cross[testIndexes, ]
      trainData <- mo_trn_cross[-testIndexes, ]
      Y_CV <- testData$price_range
      #Use the test and train data partitions however you desire...

      #num_features <- sqrt(length(colnames(mo_trn)) -1)

      rf_classifier <- randomForest(price_range ~ ., data = trainData, ntree = num_trees, mtry =
num_pred, importance = TRUE )
      Y_test_hat <- predict(rf_classifier, newdata = testData, type = "class")


      accuracy <- mean(Y_test_hat == Y_CV)

      accuracies <- c(accuracies, accuracy)
    }
    CV_accuracy <- mean(accuracies)
    CV_accuracies <- c(CV_accuracies, CV_accuracy)
  }
  #print(num_pred)
  dd <- cbind(dd, CV_accuracies)
}
dd[, num_tree := trees]
colnames(dd) <- c("m=p", "m=p/2", "m=sqrt(p)", "num_tree")
dd

```

##	m=p	m=p/2	m=sqrt(p)	num_tree
## 1:	0.8764286	0.8671429	0.7871429	10
## 2:	0.8864286	0.8807143	0.8271429	20
## 3:	0.8921429	0.8878571	0.8407143	30
## 4:	0.8957143	0.8864286	0.8428571	40
## 5:	0.8850000	0.8971429	0.8507143	50
## 6:	0.8885714	0.8914286	0.8642857	60
## 7:	0.8850000	0.8928571	0.8592857	70
## 8:	0.8800000	0.8957143	0.8607143	80
## 9:	0.8864286	0.8921429	0.8671429	90
## 10:	0.8857143	0.9007143	0.8700000	100
## 11:	0.8885714	0.9078571	0.8742857	110
## 12:	0.8828571	0.8928571	0.8735714	120
## 13:	0.8914286	0.8964286	0.8800000	130
## 14:	0.8942857	0.8935714	0.8728571	140
## 15:	0.8878571	0.9057143	0.8750000	150
## 16:	0.8978571	0.9007143	0.8735714	160
## 17:	0.8914286	0.8950000	0.8764286	170
## 18:	0.8771429	0.8950000	0.8721429	180
## 19:	0.8900000	0.8900000	0.8821429	190
## 20:	0.8921429	0.8907143	0.8742857	200
## 21:	0.8864286	0.8928571	0.8707143	210
##	m=p	m=p/2	m=sqrt(p)	num_tree

Plotting Accuracy vs Number of Trees for various values of m:



Predictions with Optimal Hyperparameters:

```

rf_classifier <- randomForest(price_range ~ ., data = mo_trn, ntree = 110, mtry = p_ov2, importance = TRUE )

Y_test_hat <- predict(rf_classifier, newdata = mo_test, type = "class")

test_accuracy <- mean(Y_test_hat == Y_test)

cm <- table(observed=Y_test, predicted=Y_test_hat)

diag = diag(cm) # number of correctly classified instances per class
rowsums = apply(cm, 1, sum) # number of instances per class
colsums = apply(cm, 2, sum) # number of predictions per class

precision = diag / colsums
recall = diag / rowsums
f1 = 2 * precision * recall / (precision + recall)

scores <- data.frame(precision, recall, f1)

test_accuracy

```

```
## [1] 0.8933333
```

```
cm
```

```

##      predicted
## observed  0   1   2   3
##      0 139   9   0   0
##      1   9 135   7   0
##      2   0  15 127  16
##      3   0   0   8 135

```

```
scores
```

```

##      precision    recall      f1
## 0 0.9391892 0.9391892 0.9391892
## 1 0.8490566 0.8940397 0.8709677
## 2 0.8943662 0.8037975 0.8466667
## 3 0.8940397 0.9440559 0.9183673

```

Variable Importance:

```

vi <- importance(rf_classifier, type = 2)

vi

```

```
##           MeanDecreaseGini
## battery_power      104.622169
## blue               1.877017
## clock_speed        12.654598
## dual_sim           2.706942
## fc                 11.300411
## four_g             2.255693
## int_memory         16.358168
## m_dep              12.747777
## mobile_wt          20.929720
## n_cores            10.594363
## pc                 12.307350
## px_height          60.875922
## px_width           58.969501
## ram                673.496143
## sc_h               13.673196
## sc_w               12.554913
## talk_time          14.371397
## three_g            2.221251
## touch_screen       2.182862
## wifi               2.446852
```